

Deep Learning

Assignment 4

Previously in 2_fullyconnected.ipynb and 3_regularization.ipynb, we trained fully connected networks to classify notMNIST (<http://yaroslavvb.blogspot.com/2011/09/notmnist-dataset.html>) characters.

The goal of this assignment is make the neural network convolutional.

```
In [1]: # These are all the modules we'll be using later. Make sure you can import them
# before proceeding further.
from __future__ import print_function
import numpy as np
import tensorflow as tf
from six.moves import cPickle as pickle
from six.moves import range
```

```
In [2]: pickle_file = 'notMNIST.pickle'

with open(pickle_file, 'rb') as f:
    save = pickle.load(f)
    train_dataset = save['train_dataset']
    train_labels = save['train_labels']
    valid_dataset = save['valid_dataset']
    valid_labels = save['valid_labels']
    test_dataset = save['test_dataset']
    test_labels = save['test_labels']
    del save # hint to help gc free up memory
    print('Training set', train_dataset.shape, train_labels.shape)
    print('Validation set', valid_dataset.shape, valid_labels.shape)
    print('Test set', test_dataset.shape, test_labels.shape)
```

```
Training set (200000, 28, 28) (200000,)
Validation set (10000, 28, 28) (10000,)
Test set (10000, 28, 28) (10000,)
```

Reformat into a TensorFlow-friendly shape:

- convolutions need the image data formatted as a cube (width by height by #channels)
- labels as float 1-hot encodings.

```
In [3]: image_size = 28
num_labels = 10
num_channels = 1 # grayscale

import numpy as np

def reformat(dataset, labels):
    dataset = dataset.reshape(
        (-1, image_size, image_size, num_channels)).astype(np.float32)
    labels = (np.arange(num_labels) == labels[:,None]).astype(np.float32)
    return dataset, labels
train_dataset, train_labels = reformat(train_dataset, train_labels)
valid_dataset, valid_labels = reformat(valid_dataset, valid_labels)
test_dataset, test_labels = reformat(test_dataset, test_labels)
print('Training set', train_dataset.shape, train_labels.shape)
print('Validation set', valid_dataset.shape, valid_labels.shape)
print('Test set', test_dataset.shape, test_labels.shape)
```

```
Training set (200000, 28, 28, 1) (200000, 10)
Validation set (10000, 28, 28, 1) (10000, 10)
Test set (10000, 28, 28, 1) (10000, 10)
```

```
In [4]: def accuracy(predictions, labels):
        return (100.0 * np.sum(np.argmax(predictions, 1) == np.argmax(labels, 1))
                / predictions.shape[0])
```

Let's build a small network with two convolutional layers, followed by one fully connected layer. Convolutional networks are more expensive computationally, so we'll limit its depth and number of fully connected nodes.

```
In [5]: batch_size = 16
        patch_size = 5
        depth = 16
        num_hidden = 64

        graph = tf.Graph()

        with graph.as_default():

            # Input data.
            tf_train_dataset = tf.placeholder(
                tf.float32, shape=(batch_size, image_size, image_size, num_channels))
            tf_train_labels = tf.placeholder(tf.float32, shape=(batch_size, num_labels))
            tf_valid_dataset = tf.constant(valid_dataset)
            tf_test_dataset = tf.constant(test_dataset)

            # Variables.
            layer1_weights = tf.Variable(tf.truncated_normal(
                [patch_size, patch_size, num_channels, depth], stddev=0.1))
            layer1_biases = tf.Variable(tf.zeros([depth]))
            layer2_weights = tf.Variable(tf.truncated_normal(
                [patch_size, patch_size, depth, depth], stddev=0.1))
            layer2_biases = tf.Variable(tf.constant(1.0, shape=[depth]))
            layer3_weights = tf.Variable(tf.truncated_normal(
                [image_size // 4 * image_size // 4 * depth, num_hidden], stddev=0.1))
            layer3_biases = tf.Variable(tf.constant(1.0, shape=[num_hidden]))
            layer4_weights = tf.Variable(tf.truncated_normal(
                [num_hidden, num_labels], stddev=0.1))
            layer4_biases = tf.Variable(tf.constant(1.0, shape=[num_labels]))

            print(layer1_weights.get_shape())
            print(layer2_weights.get_shape())
            print(layer3_weights.get_shape())
            print(layer4_weights.get_shape())

            # Model.
            def model(data):
                conv = tf.nn.conv2d(data, layer1_weights, [1, 2, 2, 1], padding='SAME')
                # print("model: conv shape: %s" % conv.get_shape())
                hidden = tf.nn.relu(conv + layer1_biases)
                # print("model: hidden shape: %s" % hidden.get_shape())
```

```
conv = tf.nn.conv2d(hidden, layer2_weights, [1, 2, 2, 1], padding='SAME')
# print("model: conv shape: %s" % conv.get_shape())
hidden = tf.nn.relu(conv + layer2_biases)
# print("model: hidden shape: %s" % hidden.get_shape())
shape = hidden.get_shape().as_list()
reshape = tf.reshape(hidden, [shape[0], shape[1] * shape[2] * shape[3]])
hidden = tf.nn.relu(tf.matmul(reshape, layer3_weights) + layer3_biases)
return tf.matmul(hidden, layer4_weights) + layer4_biases

# Training computation.
logits = model(tf_train_dataset)
loss = tf.reduce_mean(
    tf.nn.softmax_cross_entropy_with_logits(labels=tf_train_labels, logits=logits))

# Optimizer.
optimizer = tf.train.GradientDescentOptimizer(0.05).minimize(loss)

# Predictions for the training, validation, and test data.
train_prediction = tf.nn.softmax(logits)
valid_prediction = tf.nn.softmax(model(tf_valid_dataset))
test_prediction = tf.nn.softmax(model(tf_test_dataset))
```

```
(5, 5, 1, 16)
(5, 5, 16, 16)
(784, 64)
(64, 10)
```

```
In [6]: num_steps = 1001

with tf.Session(graph=graph) as session:
    tf.global_variables_initializer().run()
    print('Initialized')
    for step in range(num_steps):
        offset = (step * batch_size) % (train_labels.shape[0] - batch_size)
        batch_data = train_dataset[offset:(offset + batch_size), :, :, :]
        batch_labels = train_labels[offset:(offset + batch_size), :]
        feed_dict = {tf_train_dataset : batch_data, tf_train_labels : batch_labels}
        _, l, predictions = session.run(
            [optimizer, loss, train_prediction], feed_dict=feed_dict)
        if (step % 50 == 0):
            print('Minibatch loss at step %d: %f' % (step, l))
            print('Minibatch accuracy: %.1f%%' % accuracy(predictions, batch_labels))
            print('Validation accuracy: %.1f%%' % accuracy(
                valid_prediction.eval(), valid_labels))
    print('Test accuracy: %.1f%%' % accuracy(test_prediction.eval(), test_labels))
```

```
Initialized
Minibatch loss at step 0: 2.675679
Minibatch accuracy: 12.5%
Validation accuracy: 10.0%
Minibatch loss at step 50: 1.455389
Minibatch accuracy: 50.0%
Validation accuracy: 50.0%
Minibatch loss at step 100: 0.369873
Minibatch accuracy: 100.0%
Validation accuracy: 74.8%
Minibatch loss at step 150: 0.495924
Minibatch accuracy: 93.8%
Validation accuracy: 76.1%
Minibatch loss at step 200: 0.417204
Minibatch accuracy: 87.5%
Validation accuracy: 76.6%
Minibatch loss at step 250: 1.547781
Minibatch accuracy: 62.5%
Validation accuracy: 75.8%
Minibatch loss at step 300: 0.860488
Minibatch accuracy: 81.2%
Validation accuracy: 80.2%
```

Minibatch loss at step 350: 0.970358
Minibatch accuracy: 62.5%
Validation accuracy: 79.3%
Minibatch loss at step 400: 0.904927
Minibatch accuracy: 87.5%
Validation accuracy: 79.6%
Minibatch loss at step 450: 0.541103
Minibatch accuracy: 81.2%
Validation accuracy: 80.3%
Minibatch loss at step 500: 0.533026
Minibatch accuracy: 93.8%
Validation accuracy: 81.0%
Minibatch loss at step 550: 0.212027
Minibatch accuracy: 93.8%
Validation accuracy: 82.0%
Minibatch loss at step 600: 0.444889
Minibatch accuracy: 87.5%
Validation accuracy: 81.4%
Minibatch loss at step 650: 0.429304
Minibatch accuracy: 87.5%
Validation accuracy: 81.1%
Minibatch loss at step 700: 0.522365
Minibatch accuracy: 81.2%
Validation accuracy: 80.9%
Minibatch loss at step 750: 0.200323
Minibatch accuracy: 100.0%
Validation accuracy: 82.3%
Minibatch loss at step 800: 0.813128
Minibatch accuracy: 87.5%
Validation accuracy: 82.1%
Minibatch loss at step 850: 0.709090
Minibatch accuracy: 81.2%
Validation accuracy: 82.4%
Minibatch loss at step 900: 0.430374
Minibatch accuracy: 87.5%
Validation accuracy: 81.2%
Minibatch loss at step 950: 0.146930
Minibatch accuracy: 100.0%
Validation accuracy: 83.3%
Minibatch loss at step 1000: 0.265595
Minibatch accuracy: 87.5%

Validation accuracy: 82.0%

Test accuracy: 89.2%

Problem 1

The convolutional model above uses convolutions with stride 2 to reduce the dimensionality. Replace the strides by a max pooling operation (`nn.max_pool()`) of stride 2 and kernel size 2.

```
In [8]: batch_size = 16
        patch_size = 5
        depth = 16
        num_hidden = 64

        beta = 0.001

        graph = tf.Graph()

        with graph.as_default():

            # Input data.
            tf_train_dataset = tf.placeholder(
                tf.float32, shape=(batch_size, image_size, image_size, num_channels))
            tf_train_labels = tf.placeholder(tf.float32, shape=(batch_size, num_labels))
            tf_valid_dataset = tf.constant(valid_dataset)
            tf_test_dataset = tf.constant(test_dataset)

            # Variables.
            layer1_weights = tf.Variable(tf.truncated_normal(
                [patch_size, patch_size, num_channels, depth], stddev=0.1))
            layer1_biases = tf.Variable(tf.zeros([depth]))
            layer2_weights = tf.Variable(tf.truncated_normal(
                [patch_size, patch_size, depth, depth], stddev=0.1))
            layer2_biases = tf.Variable(tf.constant(1.0, shape=[depth]))
            layer3_weights = tf.Variable(tf.truncated_normal(
                [image_size // 4 * image_size // 4 * depth, num_hidden], stddev=0.1))
            layer3_biases = tf.Variable(tf.constant(1.0, shape=[num_hidden]))
            layer4_weights = tf.Variable(tf.truncated_normal(
                [num_hidden, num_labels], stddev=0.1))
            layer4_biases = tf.Variable(tf.constant(1.0, shape=[num_labels]))

            print(layer1_weights.get_shape())
            print(layer2_weights.get_shape())
            print(layer3_weights.get_shape())
            print(layer4_weights.get_shape())

            # Model.
            def model(data):
                conv = tf.nn.conv2d(data, layer1_weights, [1, 1, 1, 1], padding='SAME')
                # print("model: conv shape: %s" % conv.get_shape())
```



```

pool1 = tf.nn.max_pool(conv, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')
# print("model: pool shape: %s" % pool1.get_shape())
hidden = tf.nn.relu(pool1 + layer1_biases)
# print("model: hidden shape: %s" % hidden.get_shape())

conv = tf.nn.conv2d(hidden, layer2_weights, [1, 1, 1, 1], padding='SAME')
# print("model: conv shape: %s" % conv.get_shape())
pool2 = tf.nn.max_pool(conv, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')
# print("model: pool shape: %s" % pool1.get_shape())
hidden = tf.nn.relu(pool2 + layer2_biases)
# print("model: hidden shape: %s" % hidden.get_shape())

shape = hidden.get_shape().as_list()
reshape = tf.reshape(hidden, [shape[0], shape[1] * shape[2] * shape[3]])
hidden = tf.nn.relu(tf.matmul(reshape, layer3_weights) + layer3_biases)
return tf.matmul(hidden, layer4_weights) + layer4_biases

# Training computation.
logits = model(tf_train_dataset)
loss = tf.reduce_mean(
    tf.nn.softmax_cross_entropy_with_logits(labels=tf_train_labels, logits=logits))

# Optimizer.
optimizer = tf.train.GradientDescentOptimizer(0.05).minimize(loss)

# Predictions for the training, validation, and test data.
train_prediction = tf.nn.softmax(logits)
valid_prediction = tf.nn.softmax(model(tf_valid_dataset))
test_prediction = tf.nn.softmax(model(tf_test_dataset))

```

```

(5, 5, 1, 16)
(5, 5, 16, 16)
(784, 64)
(64, 10)

```

In [9]: num_steps = 1001

```
with tf.Session(graph=graph) as session:
    tf.global_variables_initializer().run()
    print('Initialized')
    for step in range(num_steps):
        offset = (step * batch_size) % (train_labels.shape[0] - batch_size)
        batch_data = train_dataset[offset:(offset + batch_size), :, :, :]
        batch_labels = train_labels[offset:(offset + batch_size), :]
        feed_dict = {tf_train_dataset : batch_data, tf_train_labels : batch_labels}
        _, l, predictions = session.run(
            [optimizer, loss, train_prediction], feed_dict=feed_dict)
        if (step % 50 == 0):
            print('Minibatch loss at step %d: %f' % (step, l))
            print('Minibatch accuracy: %.1f%%' % accuracy(predictions, batch_labels))
            print('Validation accuracy: %.1f%%' % accuracy(
                valid_prediction.eval(), valid_labels))
    print('Test accuracy: %.1f%%' % accuracy(test_prediction.eval(), test_labels))
```

Initialized

Minibatch loss at step 0: 2.769767
Minibatch accuracy: 12.5%
Validation accuracy: 10.0%
Minibatch loss at step 50: 1.557950
Minibatch accuracy: 31.2%
Validation accuracy: 50.6%
Minibatch loss at step 100: 0.632284
Minibatch accuracy: 81.2%
Validation accuracy: 72.7%
Minibatch loss at step 150: 0.465463
Minibatch accuracy: 93.8%
Validation accuracy: 76.6%
Minibatch loss at step 200: 0.394912
Minibatch accuracy: 93.8%
Validation accuracy: 77.7%
Minibatch loss at step 250: 1.370968
Minibatch accuracy: 62.5%
Validation accuracy: 76.7%
Minibatch loss at step 300: 0.739579
Minibatch accuracy: 81.2%

Validation accuracy: 79.8%
Minibatch loss at step 350: 0.864967
Minibatch accuracy: 68.8%
Validation accuracy: 78.8%
Minibatch loss at step 400: 1.093237
Minibatch accuracy: 75.0%
Validation accuracy: 78.7%
Minibatch loss at step 450: 0.614690
Minibatch accuracy: 81.2%
Validation accuracy: 80.4%
Minibatch loss at step 500: 0.382195
Minibatch accuracy: 87.5%
Validation accuracy: 81.2%
Minibatch loss at step 550: 0.194540
Minibatch accuracy: 100.0%
Validation accuracy: 82.2%
Minibatch loss at step 600: 0.495355
Minibatch accuracy: 81.2%
Validation accuracy: 81.6%
Minibatch loss at step 650: 0.460583
Minibatch accuracy: 87.5%
Validation accuracy: 82.2%
Minibatch loss at step 700: 0.581673
Minibatch accuracy: 75.0%
Validation accuracy: 81.9%
Minibatch loss at step 750: 0.219646
Minibatch accuracy: 93.8%
Validation accuracy: 82.7%
Minibatch loss at step 800: 0.698183
Minibatch accuracy: 87.5%
Validation accuracy: 82.4%
Minibatch loss at step 850: 0.661875
Minibatch accuracy: 87.5%
Validation accuracy: 82.5%
Minibatch loss at step 900: 0.350851
Minibatch accuracy: 87.5%
Validation accuracy: 81.8%
Minibatch loss at step 950: 0.106321
Minibatch accuracy: 100.0%
Validation accuracy: 83.8%
Minibatch loss at step 1000: 0.217801
Minibatch accuracy: 93.8%

Validation accuracy: 83.2%

Test accuracy: 90.3%

```
In [10]: # at first, it failed because of memory limitation issue. Migrate it to another machine  
#
```

Problem 2

Try to get the best performance you can using a convolutional net. Look for example at the classic LeNet5 (<http://yann.lecun.com/exdb/lenet/>) architecture, adding Dropout, and/or adding learning rate decay.

```
In [24]: batch_size = 16
         patch_size = 5
         depth = 16
         num_hidden = 64

         beta = 0.001

         graph = tf.Graph()

         with graph.as_default():

             # Input data.
             tf_train_dataset = tf.placeholder(
                 tf.float32, shape=(batch_size, image_size, image_size, num_channels))
             tf_train_labels = tf.placeholder(tf.float32, shape=(batch_size, num_labels))
             tf_valid_dataset = tf.constant(valid_dataset)
             tf_test_dataset = tf.constant(test_dataset)

             # Variables.
             layer1_weights = tf.Variable(tf.truncated_normal(
                 [patch_size, patch_size, num_channels, depth], stddev=0.1))
             layer1_biases = tf.Variable(tf.zeros([depth]))
             layer2_weights = tf.Variable(tf.truncated_normal(
                 [patch_size, patch_size, depth, depth], stddev=0.1))
             layer2_biases = tf.Variable(tf.constant(1.0, shape=[depth]))
             layer3_weights = tf.Variable(tf.truncated_normal(
                 [image_size // 4 * image_size // 4 * depth, num_hidden], stddev=0.1))
             layer3_biases = tf.Variable(tf.constant(1.0, shape=[num_hidden]))
             layer4_weights = tf.Variable(tf.truncated_normal(
                 [num_hidden, num_labels], stddev=0.1))
             layer4_biases = tf.Variable(tf.constant(1.0, shape=[num_labels]))

             print(layer1_weights.get_shape())
             print(layer2_weights.get_shape())
             print(layer3_weights.get_shape())
             print(layer4_weights.get_shape())

             keep_prob = tf.placeholder("float")

             # Model.
             def model(data, keep_prob=None):
```

```

conv = tf.nn.conv2d(data, layer1_weights, [1, 1, 1, 1], padding='SAME')
# print("model: conv shape: %s" % conv.get_shape())
pool1 = tf.nn.max_pool(conv, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')
# print("model: pool shape: %s" % pool1.get_shape())
hidden = tf.nn.relu(pool1 + layer1_biases)
# print("model: hidden shape: %s" % hidden.get_shape())

conv = tf.nn.conv2d(hidden, layer2_weights, [1, 1, 1, 1], padding='SAME')
# print("model: conv shape: %s" % conv.get_shape())
pool2 = tf.nn.max_pool(conv, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')
# print("model: pool shape: %s" % pool1.get_shape())
hidden = tf.nn.relu(pool2 + layer2_biases)
# print("model: hidden shape: %s" % hidden.get_shape())

shape = hidden.get_shape().as_list()
reshape = tf.reshape(hidden, [shape[0], shape[1] * shape[2] * shape[3]])
hidden = tf.nn.relu(tf.matmul(reshape, layer3_weights) + layer3_biases)
if keep_prob is not None:
    drop = tf.nn.dropout(hidden, keep_prob)
    return tf.matmul(drop, layer4_weights) + layer4_biases
else:
    return tf.matmul(hidden, layer4_weights) + layer4_biases

# Training computation.
logits = model(tf_train_dataset, keep_prob)
loss = tf.reduce_mean(
    tf.nn.softmax_cross_entropy_with_logits(labels=tf_train_labels, logits=logits))

# Optimizer.
optimizer = tf.train.GradientDescentOptimizer(0.05).minimize(loss)

# Predictions for the training, validation, and test data.
train_prediction = tf.nn.softmax(logits)
valid_prediction = tf.nn.softmax(model(tf_valid_dataset))
test_prediction = tf.nn.softmax(model(tf_test_dataset))

```

```

(5, 5, 1, 16)
(5, 5, 16, 16)
(784, 64)
(64, 10)

```

In [25]: num_steps = 1001

```
with tf.Session(graph=graph) as session:
    tf.global_variables_initializer().run()
    print('Initialized')
    for step in range(num_steps):
        offset = (step * batch_size) % (train_labels.shape[0] - batch_size)
        batch_data = train_dataset[offset:(offset + batch_size), :, :, :]
        batch_labels = train_labels[offset:(offset + batch_size), :]
        feed_dict = {tf_train_dataset : batch_data, tf_train_labels : batch_labels, keep_prob: 0.9}
        _, l, predictions = session.run(
            [optimizer, loss, train_prediction], feed_dict=feed_dict)
        if (step % 50 == 0):
            print('Minibatch loss at step %d: %f' % (step, l))
            print('Minibatch accuracy: %.1f%%' % accuracy(predictions, batch_labels))
            print('Validation accuracy: %.1f%%' % accuracy(
                valid_prediction.eval(), valid_labels))
    print('Test accuracy: %.1f%%' % accuracy(test_prediction.eval(), test_labels))
```

```
Initialized
Minibatch loss at step 0: 4.257010
Minibatch accuracy: 6.2%
Validation accuracy: 10.0%
Minibatch loss at step 50: 2.059375
Minibatch accuracy: 18.8%
Validation accuracy: 32.8%
Minibatch loss at step 100: 1.121794
Minibatch accuracy: 62.5%
Validation accuracy: 63.2%
Minibatch loss at step 150: 0.840861
Minibatch accuracy: 68.8%
Validation accuracy: 72.6%
Minibatch loss at step 200: 0.615932
Minibatch accuracy: 87.5%
Validation accuracy: 76.9%
Minibatch loss at step 250: 1.432374
Minibatch accuracy: 68.8%
Validation accuracy: 75.2%
Minibatch loss at step 300: 0.574929
Minibatch accuracy: 87.5%
```

Validation accuracy: 79.8%
Minibatch loss at step 350: 1.029322
Minibatch accuracy: 56.2%
Validation accuracy: 79.7%
Minibatch loss at step 400: 1.073759
Minibatch accuracy: 81.2%
Validation accuracy: 79.4%
Minibatch loss at step 450: 0.679190
Minibatch accuracy: 81.2%
Validation accuracy: 81.0%
Minibatch loss at step 500: 0.542995
Minibatch accuracy: 81.2%
Validation accuracy: 82.2%
Minibatch loss at step 550: 0.264706
Minibatch accuracy: 93.8%
Validation accuracy: 81.9%
Minibatch loss at step 600: 0.479310
Minibatch accuracy: 87.5%
Validation accuracy: 82.4%
Minibatch loss at step 650: 0.530982
Minibatch accuracy: 87.5%
Validation accuracy: 82.4%
Minibatch loss at step 700: 0.582734
Minibatch accuracy: 81.2%
Validation accuracy: 82.5%
Minibatch loss at step 750: 0.252468
Minibatch accuracy: 100.0%
Validation accuracy: 83.0%
Minibatch loss at step 800: 0.857202
Minibatch accuracy: 87.5%
Validation accuracy: 83.1%
Minibatch loss at step 850: 0.679520
Minibatch accuracy: 75.0%
Validation accuracy: 82.4%
Minibatch loss at step 900: 0.457159
Minibatch accuracy: 81.2%
Validation accuracy: 82.2%
Minibatch loss at step 950: 0.201450
Minibatch accuracy: 100.0%
Validation accuracy: 83.8%
Minibatch loss at step 1000: 0.284401
Minibatch accuracy: 87.5%

Validation accuracy: 82.8%

Test accuracy: 89.2%

In [26]: *# it seems the starter Learning rate is a little high. Then at the end, it cannot get a better result.*

```
In [29]: batch_size = 16
patch_size = 5
depth = 16
num_hidden = 64
beta = 0.01

graph = tf.Graph()

with graph.as_default():

    # Input data.
    tf_train_dataset = tf.placeholder(
        tf.float32, shape=(batch_size, image_size, image_size, num_channels))
    tf_train_labels = tf.placeholder(tf.float32, shape=(batch_size, num_labels))
    tf_valid_dataset = tf.constant(valid_dataset)
    tf_test_dataset = tf.constant(test_dataset)

    # Variables.
    layer1_weights = tf.Variable(tf.truncated_normal(
        [patch_size, patch_size, num_channels, depth], stddev=0.1))
    layer1_biases = tf.Variable(tf.zeros([depth]))
    layer2_weights = tf.Variable(tf.truncated_normal(
        [patch_size, patch_size, depth, depth], stddev=0.1))
    layer2_biases = tf.Variable(tf.constant(1.0, shape=[depth]))
    layer3_weights = tf.Variable(tf.truncated_normal(
        [image_size // 4 * image_size // 4 * depth, num_hidden], stddev=0.1))
    layer3_biases = tf.Variable(tf.constant(1.0, shape=[num_hidden]))
    layer4_weights = tf.Variable(tf.truncated_normal(
        [num_hidden, num_labels], stddev=0.1))
    layer4_biases = tf.Variable(tf.constant(1.0, shape=[num_labels]))

    print(layer1_weights.get_shape())
    print(layer2_weights.get_shape())
    print(layer3_weights.get_shape())
    print(layer4_weights.get_shape())

    # Model.
    def model(data, keep_prob=None):
        conv = tf.nn.conv2d(data, layer1_weights, [1, 1, 1, 1], padding='SAME')
        # print("model: conv shape: %s" % conv.get_shape())
        pool1 = tf.nn.max_pool(conv, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')
```

```

# print("model: pool shape: %s" % pool1.get_shape())
hidden = tf.nn.relu(pool1 + layer1_biases)
# print("model: hidden shape: %s" % hidden.get_shape())

conv = tf.nn.conv2d(hidden, layer2_weights, [1, 1, 1, 1], padding='SAME')
# print("model: conv shape: %s" % conv.get_shape())
pool2 = tf.nn.max_pool(conv, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')
# print("model: pool shape: %s" % pool1.get_shape())
hidden = tf.nn.relu(pool2 + layer2_biases)
# print("model: hidden shape: %s" % hidden.get_shape())

shape = hidden.get_shape().as_list()
reshape = tf.reshape(hidden, [shape[0], shape[1] * shape[2] * shape[3]])
hidden = tf.nn.relu(tf.matmul(reshape, layer3_weights) + layer3_biases)
if keep_prob is not None:
    drop = tf.nn.dropout(hidden, keep_prob)
    return tf.matmul(drop, layer4_weights) + layer4_biases
else:
    return tf.matmul(hidden, layer4_weights) + layer4_biases

keep_prob = tf.placeholder("float")
# Training computation.
logits = model(tf_train_dataset, keep_prob)
loss = tf.reduce_mean(
    tf.nn.softmax_cross_entropy_with_logits(labels=tf_train_labels, logits=logits)\
    + beta * tf.nn.l2_loss(layer4_weights)\
    + beta * tf.nn.l2_loss(layer3_weights)\
    + beta * tf.nn.l2_loss(layer2_weights)\
    + beta * tf.nn.l2_loss(layer1_weights))

# Optimizer.
#optimizer = tf.train.GradientDescentOptimizer(0.05).minimize(loss)
global_step = tf.Variable(0) # count the number of steps taken.
starter_learning_rate = tf.placeholder("float")
learning_rate = tf.train.exponential_decay(starter_learning_rate, global_step, 10000, 0.96, staircase=True)
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss, global_step=global_step)

# Predictions for the training, validation, and test data.
train_prediction = tf.nn.softmax(logits)
valid_prediction = tf.nn.softmax(model(tf_valid_dataset))
test_prediction = tf.nn.softmax(model(tf_test_dataset))

```

(5, 5, 1, 16)
(5, 5, 16, 16)
(784, 64)
(64, 10)



```

In [34]: num_steps = 3001
         keep_prob_rate = 0.9
         starter_learning = 0.00001

         with tf.Session(graph=graph) as session:
             tf.global_variables_initializer().run()
             print('Initialized')
             for step in range(num_steps):
                 offset = (step * batch_size) % (train_labels.shape[0] - batch_size)
                 batch_data = train_dataset[offset:(offset + batch_size), :, :, :]
                 batch_labels = train_labels[offset:(offset + batch_size), :]
                 feed_dict = {tf_train_dataset : batch_data, tf_train_labels : batch_labels, keep_prob: 0.9, starter_learning_rate: 0.
                 _, l, predictions = session.run(
                     [optimizer, loss, train_prediction], feed_dict=feed_dict)
                 if (step % 50 == 0):
                     print('Minibatch loss at step %d: %f' % (step, l))
                     print('Minibatch accuracy: %.1f%%' % accuracy(predictions, batch_labels))
                     print('Validation accuracy: %.1f%%' % accuracy(
                         valid_prediction.eval(), valid_labels))
                     print('Test accuracy: %.1f%%' % accuracy(test_prediction.eval(), test_labels))
Minibatch accuracy: 88.8%
Validation accuracy: 80.6%
Minibatch loss at step 2750: 2.721976
Minibatch accuracy: 75.0%
Validation accuracy: 80.5%
Minibatch loss at step 2800: 2.636328
Minibatch accuracy: 81.2%
Validation accuracy: 80.6%
Minibatch loss at step 2850: 2.731177
Minibatch accuracy: 81.2%
Validation accuracy: 80.7%
Minibatch loss at step 2900: 2.495913
Minibatch accuracy: 87.5%
Validation accuracy: 80.8%
Minibatch loss at step 2950: 2.431013
Minibatch accuracy: 81.2%
Validation accuracy: 80.7%
Minibatch loss at step 3000: 2.240144
Minibatch accuracy: 87.5%
Validation accuracy: 80.7%

```

In []: